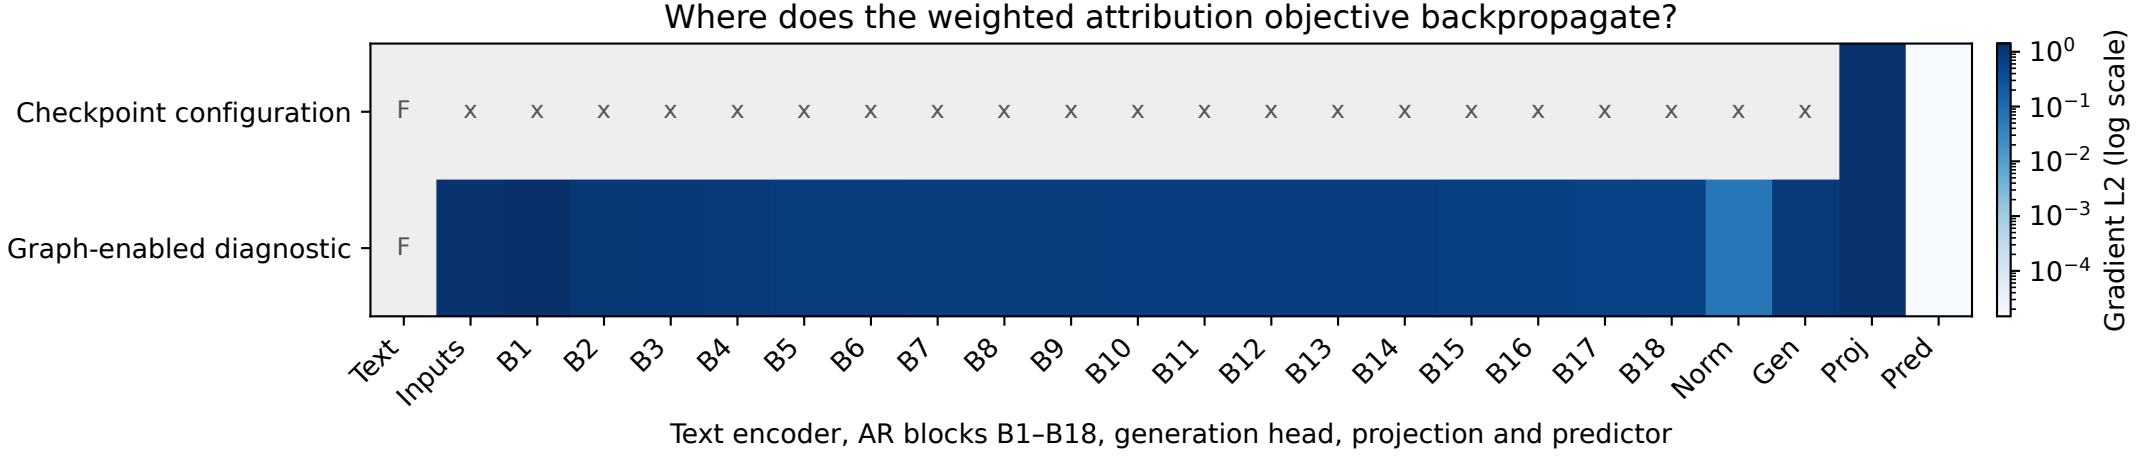
Where does the weighted attribution objective backpropagate?
Checkpoint configuration
Graph-enabled diagnostic
F
x
Text
Inputs
B1
B2
B3
B4
B5
B6
B7
B8
B9
B10
B11
B12
B13
B14
B15
B16
B17
B18
Norm
Gen
Proj
Pred
Text encoder, AR blocks B1–B18, generation head, projection and predictor
Gradient L2 (log scale)
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}